KPLABS Course

Certified Kubernetes Security Specialist

System Hardening

ISSUED BY
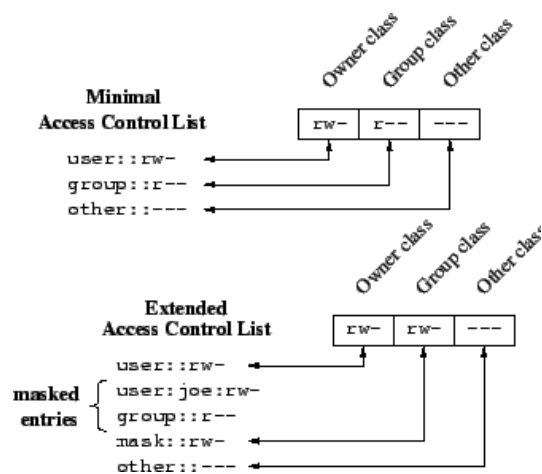
Zeal

REPRESENTATIVE

instructors@kplabs.in

Module 1: Overview of AppArmor

1.1 Revising DAC

Discretionary access control (DAC) allows restricting access to objects based on the identity of subjects and/or groups to which they belong.

```
                                          Owner class   Group class   Other class
        Minimal
     Access Control List              |  rw-  |  r--  |  ---  |
        user::rw-    ◄
        group::r--   ◄
        other::---   ◄
                                          Owner class   Group class   Other class
        Extended
     Access Control List              |  rw-  |  rw-  |  ---  |
        user::rw-    ◄
  masked  ┌ user:joe:rw-
  entries └ group::r--
        mask::rw-    ◄
        other::---   ◄
```

1.2 Challenges with DAC

Many of the times, the permissions on the files are more than what is needed.

The programs that are running with the user privilege has complete access to the data owned by that user.

1.3 Sample Use-Case

A client has sent you a zip file on the pretext that it contains certain images related to the error that he is facing while using your program.

In the background, the zip file takes some sensitive data from your host and transmits to the internet.

```
┌──────────┐   Uploading data to Server   ┌──────────┐
│ ZIP File │   ───────────────────────▶   │  Server  │
└──────────┘                              └──────────┘
```

1.4 Mandatory Access Control

Mandatory Access Control is set by the Administrator.

Two important concepts: Confined (Not Trusted) and Unconfined (Trusted)

```
┌──────────────┐          ┌──────────────┐
│  Process X   │          │  Process Z   │
└──────────────┘          └──────────────┘
   Confined                 Not  Confined
```

1.5 Confined Process

Confined Processes are not trusted.

Everything that the process intends to do must be listed in a profile.

If that capability is not listed in the profile, the process will not be allowed to run that.

```
┌──────────┐                    ┌──────────────┐
│ ZIP File │   ◀─────────▶      │  Read Files  │
└──────────┘                    └──────────────┘
```

Following are the two primary modes for AppArmor

Modes	Description
Enforce	Enforces the policy defined in the profile.
Complain	In complain or learning mode, violations of AppArmor profile rules, such as the profiled program accessing files not permitted by the profile, are detected. The violations are permitted, but also logged

Module 2: Integrate AppArmor with K8s

Kubernetes support for Apparmor was added in v1.4

AppArmor profiles are specified per-container.

To specify the AppArmor profile to run a Pod container with, add an annotation to the Pod's metadata:

```
container.apparmor.security.beta.kubernetes.io/nginx: localhost/root.apparmor.myscript
```

It is important to ensure that an apparmor profile is available in the node where the pod will be scheduled.

```
root@ip-172-31-59-221:~# kubectl get pods
NAME                READY    STATUS     RESTARTS    AGE
hello-apparmor      0/1      Blocked    0           24s
```

Module 3: OCI and Container Runtimes

3.1 Importance of Standardization

In an organization, if image standardization is not set, different developers will use a different set of images.

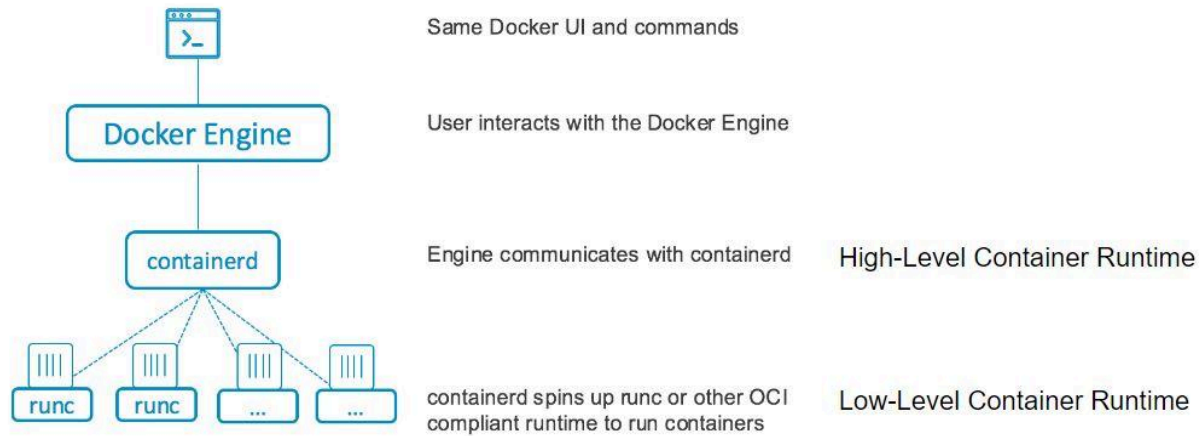This leads to challenges in troubleshooting, as well as security.

3.2 Open Container Initiative

The Open Container Initiative (OCI) is a Linux Foundation project to design open standards for containers.

There are two important specifications

Specification	Description
Image Specification	Defines how to create an OCI Image, which includes an image manifest, a filesystem (layer) serialization, and an image configuration.
Runtime Specification	defines how to run the OCI image bundle as a container.

3.3 Docker Workflow

Same Docker UI and commands

Docker Engine — User interacts with the Docker Engine

containerd — Engine communicates with containerd — High-Level Container Runtime

runc, runc, ..., ... — containerd spins up runc or other OCI compliant runtime to run containers — Low-Level Container Runtime

3.4 Container Runtimes

A container runtime is a software that executes containers and manages container images on a node

There are multiple container runtimes available. Some of these include:

- Docker
- containerd
- Cri-o
- Podman

- Pulling Images from registry.
- Unpacks image into containers root fs
- Generates OCI runtime spec json describing how to run container
- Launches OCI compatible runtime (default runc)

- Runs the container process
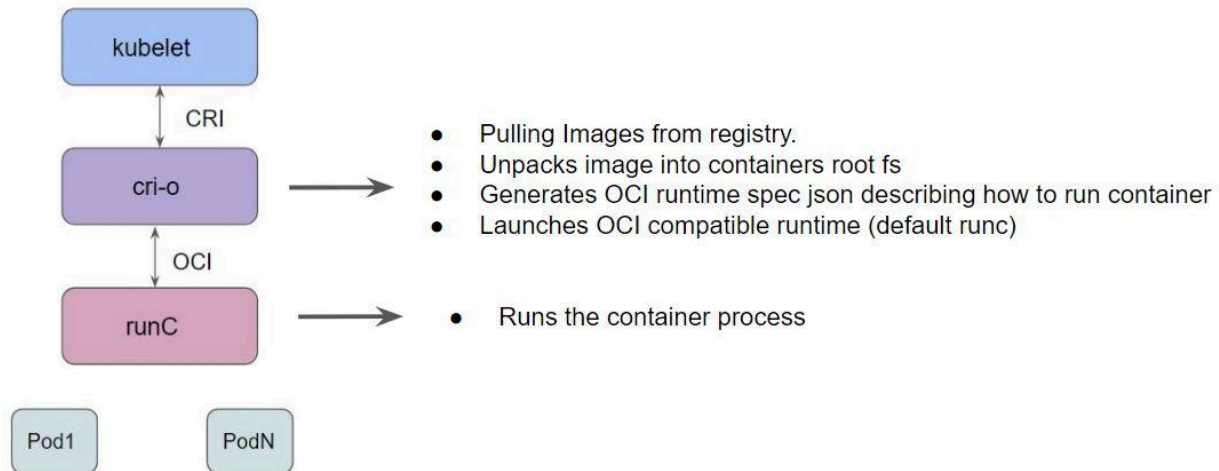
Module 4: Container Runtime - Implementation

We will focus on two implementation approaches:

1. Creating containers via containerd
2. Creating containers directly from runC

Module 5: Container Runtime Interface

Each container runtime has its own strengths.

K8s uses Container Runtime Interface which is a plugin interface that enables kubelet to use a wide variety of container runtimes without the need to recompile.

The following diagram shows a sample architecture on how it integrates.

.
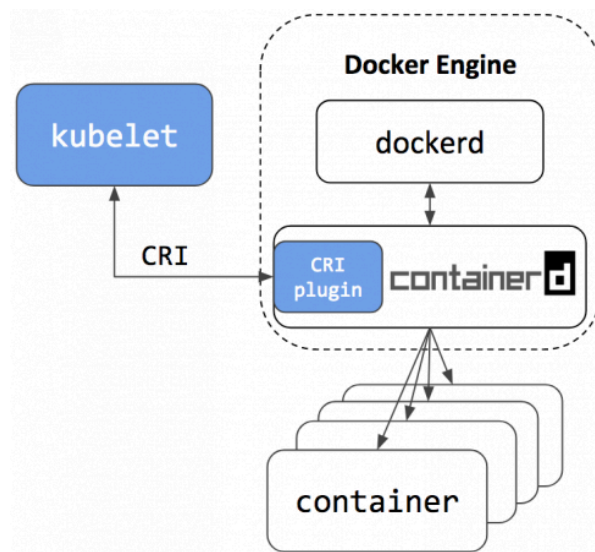
Module 6: Container Sandbox

6.1 Basic Architecture

Applications that run in traditional Linux containers access system resources in the same way that regular (non-containerized) applications do: by making system calls directly to the host kernel.

6.2 Use Cases - Bugs

If there are certain bugs at the Kernel level, the application can take advantage of it to achieve various use-cases like privilege escalation and others.

6.3 Seccomp

Kernel features like seccomp filters can provide better isolation between the application and host kernel, but they require the user to create a predefined whitelist of system calls.

In practice, it's often difficult to know which system calls will be required by an application beforehand.


6.4 Container Sandbox

Sandboxing is an approach that enforces a level of isolation between the software running on the machine and the underlying operating system.

In practice, it's often difficult to know which system calls will be required by an application beforehand.

6.5 Overview of gVisor

The core of gVisor is a kernel that runs as a normal, unprivileged process that supports most Linux system calls.

gVisor intercepts all system calls made by the application and does the necessary work to service them thus providing a strong isolation boundary.

6.6 Exploring gVisor

It primarily replaces runc (default runtime) which had few serious vulnerabilities

It comes with an OCI runtime named runsc and hence can act as a drop-in replacement to the runc.

6.7 Exploring dmesg

dmesg (diagnostic message) is a command on most Unix-like operating systems that prints the message buffer of the kernel.

gVisor based POD

Default runtime class pod

6.8 Challenges

Generally, an organization makes use of sandboxes like gVisor for the applications that are not entirely trusted (cloning repo from GitHub and running that application)

It can lead to certain performance degradation.

Module 7: RunTimeClass

RuntimeClass is a feature for selecting the container runtime configuration.

You can set a different RuntimeClass between different Pods to provide a balance of performance versus security.


```
apiVersion: node.k8s.io/v1
kind: RuntimeClass
metadata:
  name: gvisor
handler: runsc
```

Module 8: Kubeadm with Calico

8.1 Overview of Flannel

Flannel is one of the projects that is developed by CoreOS and is one of the most straightforward and popular CNI plugins.

We have been using Flannel for most of our demos and use-cases.

Flannel is generally considered as the default simple choice however Calico is also popular because of its flexibility and performance.

Calico is also known for its various advanced network capabilities and Network Policy is one of the popular capabilities that is sought.


Module 9: Network Policies In-Detail

9.1 Basics of Network Policies

By default, pods are non-isolated; they accept traffic from any source.

A network policy is a specification of how groups of pods are allowed to communicate with each other and other network endpoints.

Example: POD 1 should not have any internal or external communication.

.

9.2 Default Deny - Ingress

This will block all the inbound traffic to the pods in the namespace.

Egress traffic will still be allowed to pass.

```yaml
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-ingress
spec:
  podSelector: {}
  policyTypes:
  - Ingress
```

9.3 Example Scenario - Deny Ingress

Since all ingress is blocked, the following things will happen:

- New incoming connection to the PODS would not be allowed.
- PODS would be allowed to communicate externally and return traffic will be allowed.

.

9.4 Default Deny - Egress

This will block all the outbound traffic to the pods.

Ingress traffic will still be allowed.

```yaml
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: default-deny-egress
spec:
  podSelector: {}
  policyTypes:
  - Egress
```

9.5 Example Scenario - Deny Egress

Since all egress is blocked, the following things will happen:

- Outgoing connections from the POD will not be allowed.
- Incoming Connections will be allowed.

9.6 POD Selector

Allows us to group the pods to which the policy should apply.

All Ingress and Egress traffic for POD 1 should be blocked.


```yaml
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: deny-ingress-egress
spec:
  podSelector:
    matchLabels:
      role: suspicious
  policyTypes:
  - Ingress
  - Egress
```

9.7 Ingress - From Selector

Each rule allows traffic that matches both the from and ports sections

Rules
Allow from 192.168.0.0/16
Deny from 192.168.188.5

```yaml
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: ingress-from-ips
spec:
  podSelector:
    matchLabels:
      role: secure
  ingress:
  - from:
    - ipBlock:
        cidr: 192.168.0.0/16
        except:
        - 192.168.182.197/32
  policyTypes:
  - Ingress
```

9.8 Egress - To Selector

Allows outgoing connections only to a certain set of networks.


```yaml
---
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: egress-to-ips
spec:
  podSelector:
    matchLabels:
      role: secure
  egress:
  - to:
    - ipBlock:
        cidr: 192.168.182.197/32
  policyTypes:
  - Egress
```

9.9 Namespace Selector

Allow connections based on namespace.

Example:

Allow ingress connections to FTP POD from namespace app with pods labeled reconcile.

Join Our Discord Community

We invite you to join our Discord community, where you can interact with our support team for any course-based technical queries and connect with other students who are doing the same course.

Joining URL:

http://kplabs.in/chat